

# Introduction to Deep Learning with Tensorflow and Keras libraries with some examples in biomedical research at the Hospital Clínic of Barcelona

Roger Borràs

SEA

May 2018



## 1 INTRODUCTION

- Amazing examples
- Motivation

## 2 NEURAL NETWORKS

- Architectures
- Optimization, backpropagation & gradient descent
- Tips & Tricks

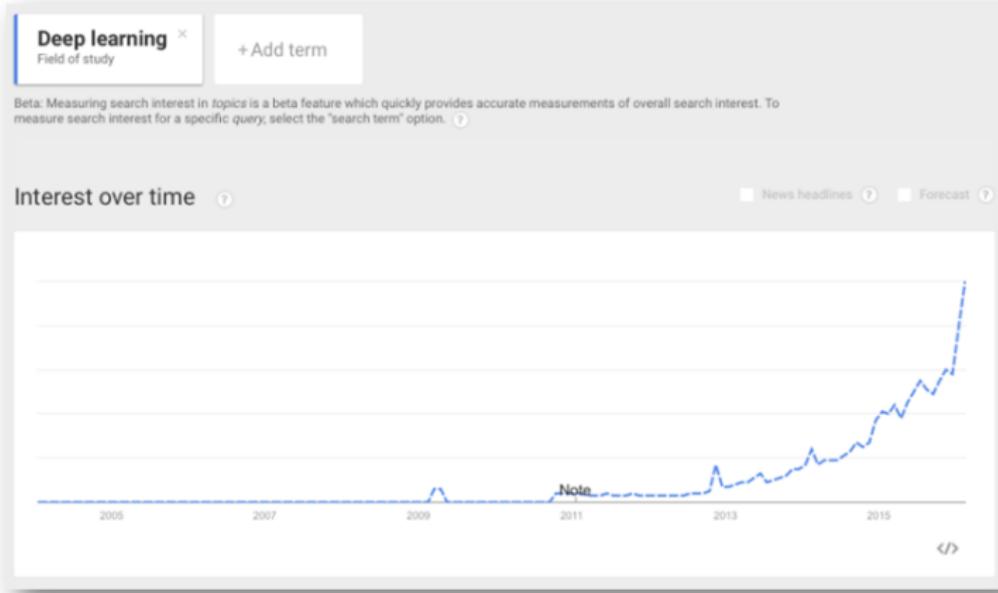
## 3 EXAMPLES WITH TensorFlow and Keras

- Tools for deep learning
- Imperative vs. Symbolic API
- Example: Linear regression, Logistic regression and FCL
- Example: CNN
- Example: LSTM
- Advanced features (more examples) and comments

## 4 CONCLUSIONS

# INTRODUCTION

# Interest



# Functions a Deep Neural Network can learn

## input

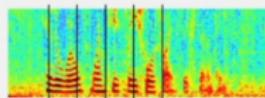
Pixels:



## output

"lion"

Audio:



"How cold is it outside?"

"Hello, how are you?"

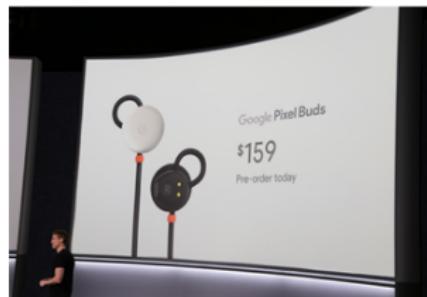
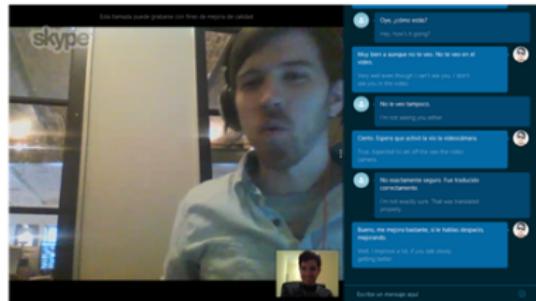
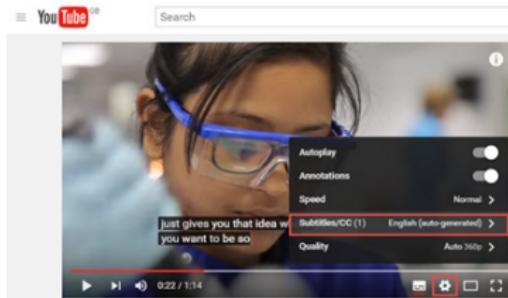
"Bonjour, comment allez-vous?"

Pixels:



"A blue and yellow train travelling down the tracks"

# Functions a Deep Neural Network can learn



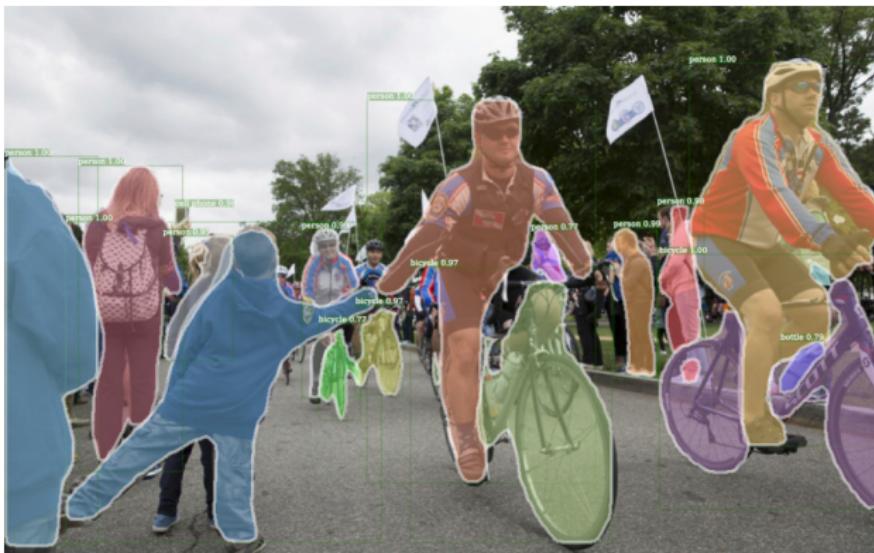
# Functions a Deep Neural Network can learn



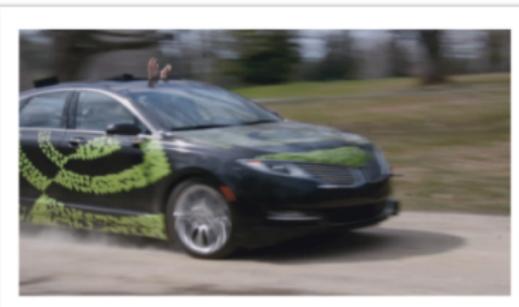
# Functions a Deep Neural Network can learn



# Functions a Deep Neural Network can learn



# Functions a Deep Neural Network can learn



# Question answering model

## Conversation 4: Philosophical Q&A

**Human:** what is the purpose of life ?

**Machine:** to serve the greater good .

**Human:** what is the purpose of living ?

**Machine:** to live forever .

**Human:** what is the purpose of existence ?

**Machine:** to find out what happens when we get to the planet earth .

**Human:** where are you now ?

**Machine:** i 'm in the middle of nowhere .

**Human:** what is the purpose of dying ?

**Machine:** to have a life .

**Human:** what is the purpose of being intelligent ?

**Machine:** to find out what it is .

**Human:** what is the purpose of emotions ?

**Machine:** i don 't know .

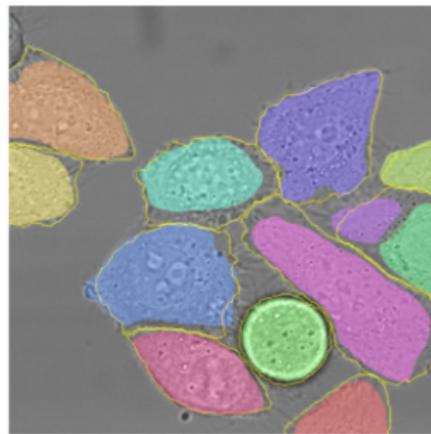
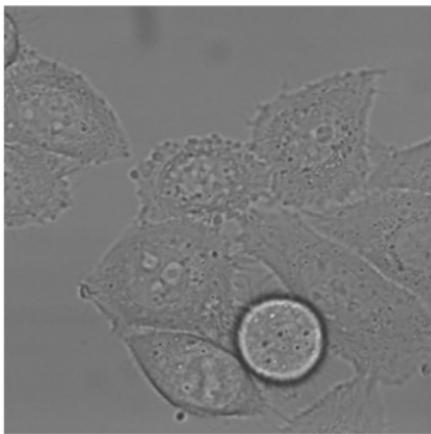
# Deep adversarial learning



# Neural Style Transfer



# Healthcare Applications

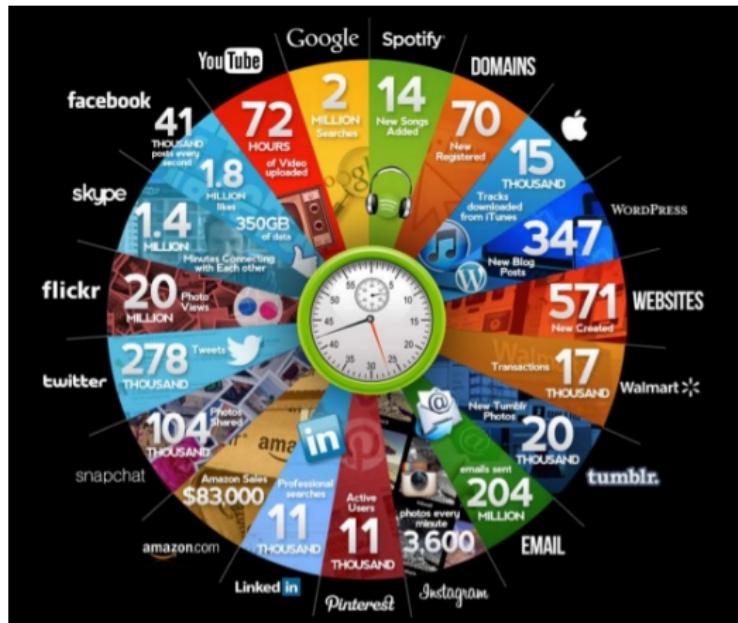


# Error in automated Image Captioning



"a young boy is holding a  
baseball bat."

# Machine Learning perspective



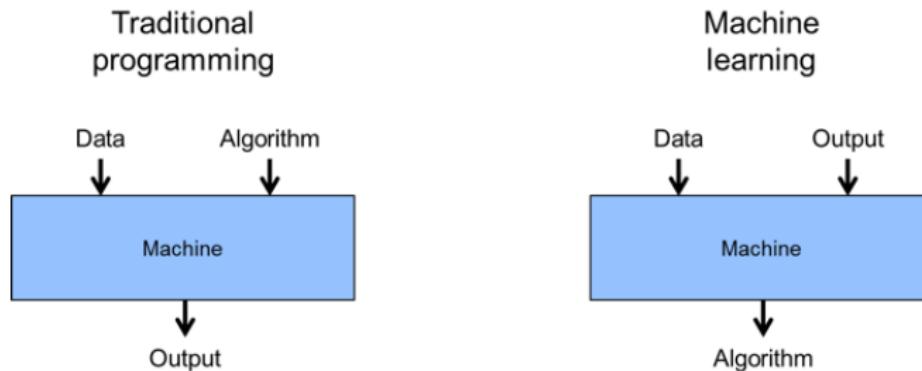
# Machine Learning: The concept

## Types of Machine Learning

Machine  
Learning



# Machine Learning perspective



# Software 2.0 with Deep Neural Networks

## Medium



Andrej Karpathy [Follow](#)

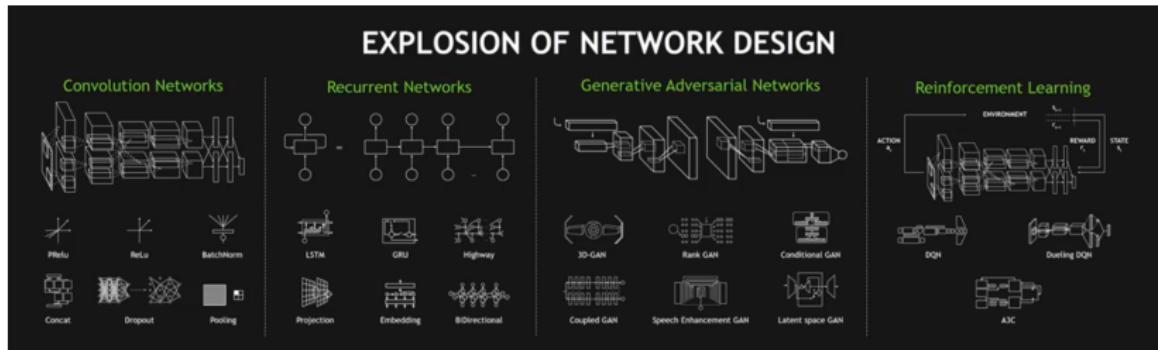
Director of AI at Tesla. Previously Research Scientist at OpenAI and PhD student at Stanford. I like to train deep neural nets on large datasets.

Nov 12, 2017 · 8 min read

## Software 2.0

<https://medium.com/@karpathy/software-2-0-a64152b37c35>

# Networks designs



# Functions a Deep Neural Network can learn

Cats vs. Dogs image Classification...



Sample of cats & dogs images from Kaggle Dataset

# Magic? AI?

# Properties

 François Chollet   
@fchollet

Follow ▾

"Neural networks" are a sad misnomer.  
They're neither neural nor even  
networks. They're chains of  
differentiable, parameterized geometric  
functions, trained with gradient descent  
(with gradients obtained via the chain  
rule). A small set of highschool-level  
ideas put together

11:58 AM - 12 Jan 2018

1,319 Retweets 3,423 Likes

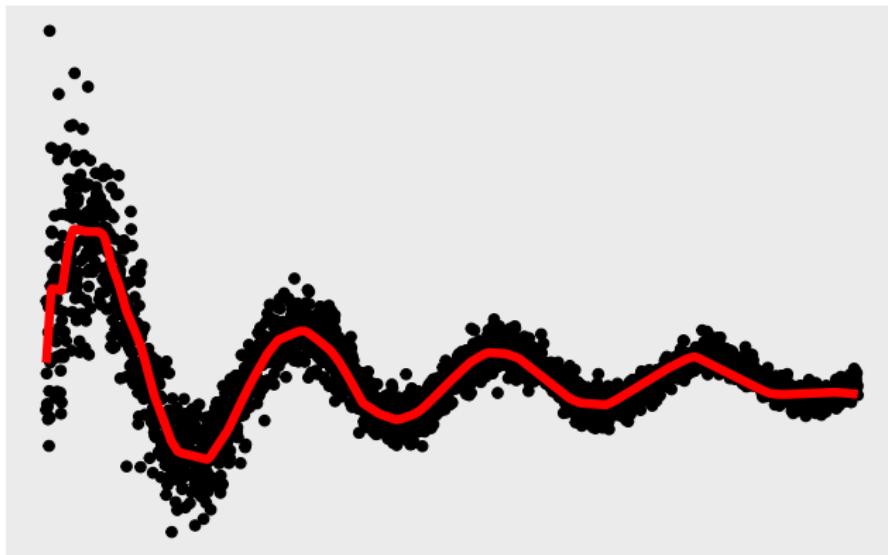


114 1.3K 3.4K

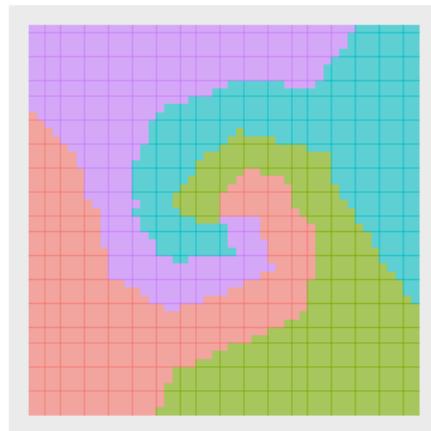
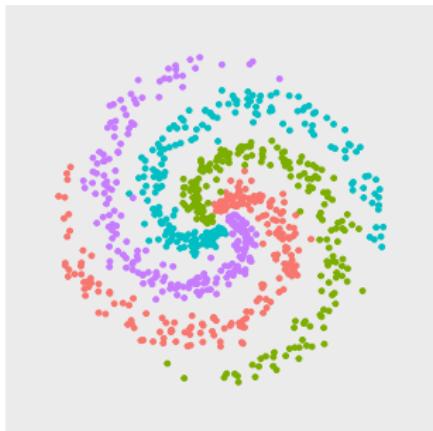
A simple mechanism that, once scaled, ends up looking like magic.

## Regression problem: $(1/x) \sin(x)$

Powerful models to map inputs to outputs!

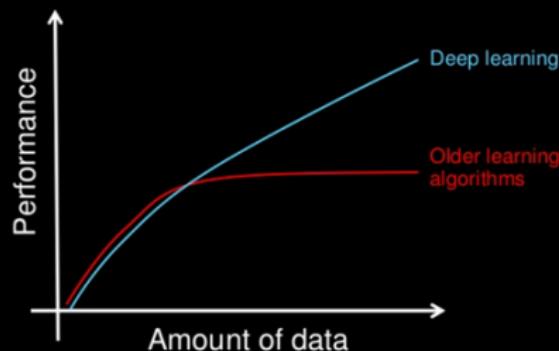


# Classification problem: Spiral data



# Properties

## Why deep learning

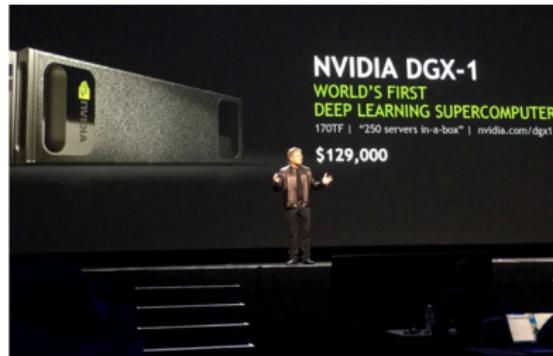


How do data science techniques scale with amount of data?

# GPU computing



NVIDIA TITAN X  
THE ULTIMATE.



Deep Learning performance doubles  
For data scientists and researchers

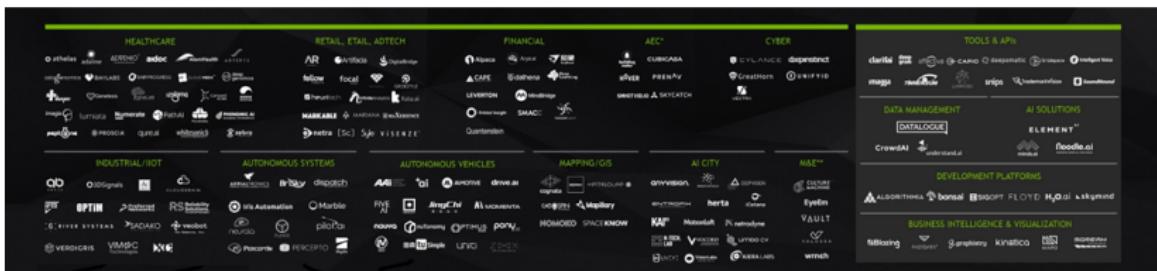


2x Faster Single GPU Training  
Support for Larger Models



2x Larger Datasets  
Instruction-level Profiling

# Companies and startups that use DL



# NEURAL NETWORKS

## General formulation

In the function estimation problem one has a system consisting of a random output variable  $\mathbf{y}$  and a set of random input variables  $\mathbf{x} = \{x_1, x_2, \dots, x_n\}$ . Given a *training* sample  $\{y_i, x_i\}_1^N$ , the goal is to find a function  $F_\theta^*(\mathbf{x})$  that maps  $\mathbf{x}$  to  $\mathbf{y}$ , such that over the joint distribution of all  $(\mathbf{y}, \mathbf{x})$  values, the expected value of some specified loss function  $J(y, F_\theta(\mathbf{x}))$  is minimized given  $\theta$ :

$$F_\theta^*(\mathbf{x}) = \operatorname{argmin}_{\theta} J(y, F_\theta(\mathbf{x}))$$
$$\theta \in \mathbb{R}^n$$

## Example: linear regression

A multiple **linear regression** allows to express the relationship between inputs and outputs as follows:

$$Y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_n X_{in}$$

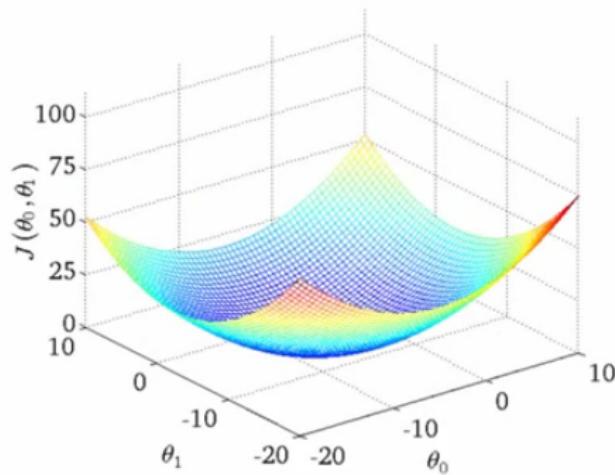
The loss function is minimizing SE:

$$J(\beta) = \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \dots + \beta_n X_{in}))^2$$

In this particular case, we can obtain a closed formula for the optimal parameters:

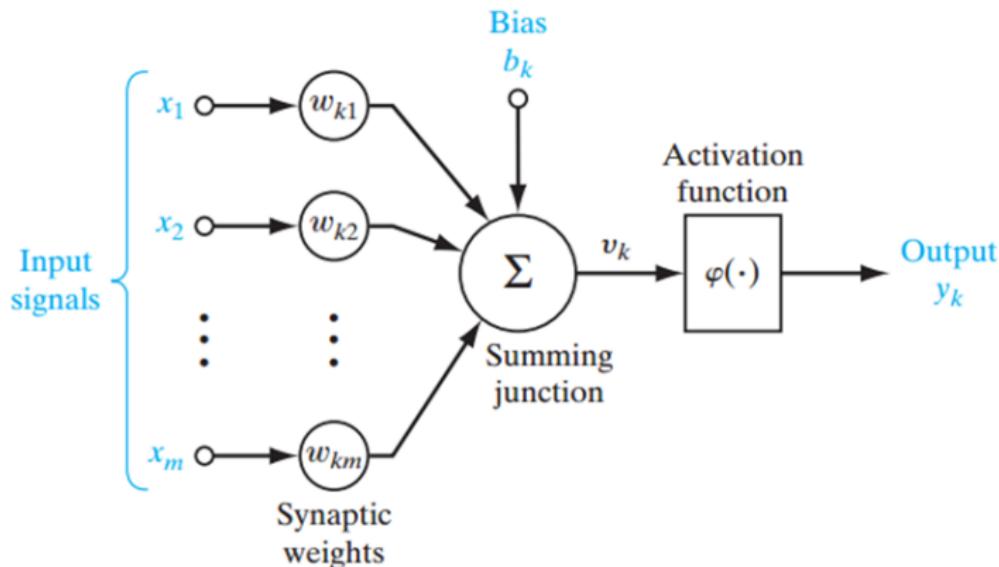
$$\hat{\beta} = (X^T X)^{-1} (X^T Y)$$

## Loss function in LR



Convex optimization problem (quadratic form)

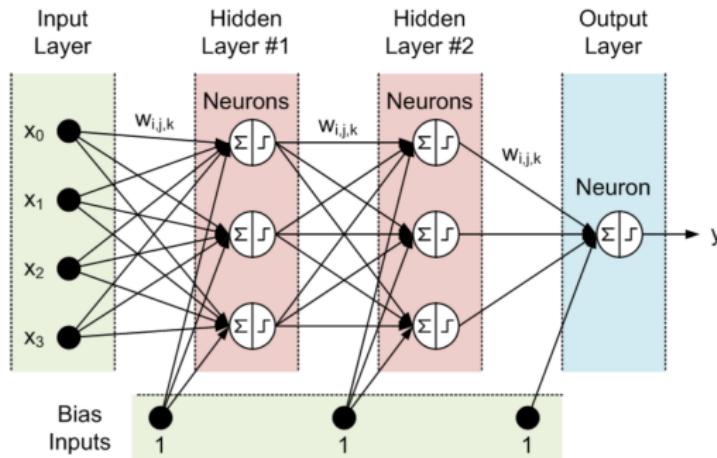
# Neuron



# Why use hidden layers configuration?



## Why use hidden layers configuration?

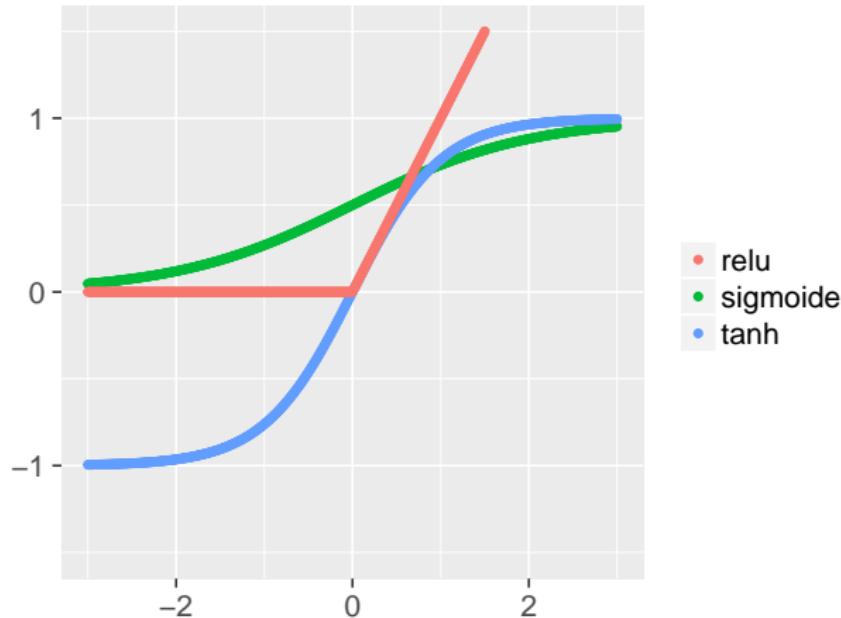


The universal approximation theorem found that a neural network with one hidden layer can approximate any continuous function (George Cybenko, 1989).

# Activation functions

Name	Visualization	$f(x) =$	Notes
Linear (= Identity)		$x$	Not useful for hidden layers
Heaviside Step		$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	Not differentiable
Rectified Linear (ReLU)		$\begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	Surprisingly useful in practice
Tanh		$\frac{e^x - e^{-x}}{e^x + e^{-x}} - 1$	A soft step function; ranges from -1 to 1
Logistic ('sigmoid')		$\frac{1}{1+e^{-x}}$	Another soft step function; ranges from 0 to 1

# Activation functions



# Loss functions

Error Metrics for Regression Problems:

- Mean Squared Error:  $(\frac{1}{n} \sum (y_i - \hat{y}_i)^2)$ .
- Root Mean Squared Error.
- ...

Error Metrics for Classification Problems:

- Logistic Loss.
- Logarithmic Loss:  $-\frac{1}{N} \sum \sum y_{ij} \log(p_{ij})$ .
- Accuracy.
- ...

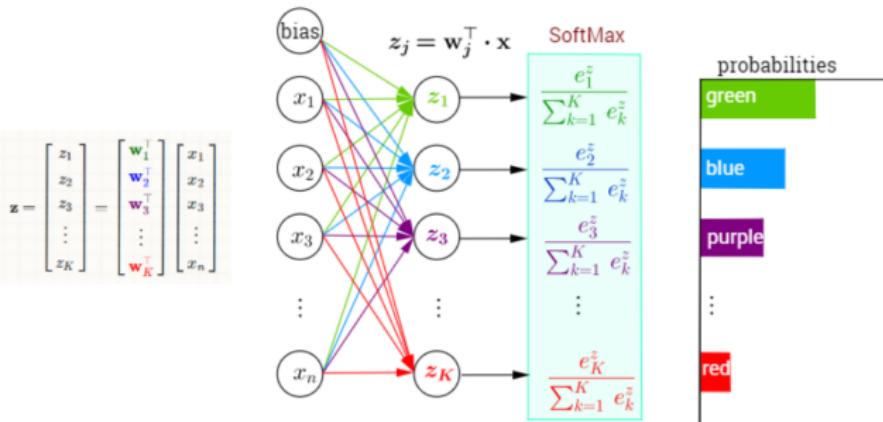
## Softmax function activation

$$\sigma : \mathbb{R}^K \rightarrow [0, 1]^K$$
$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

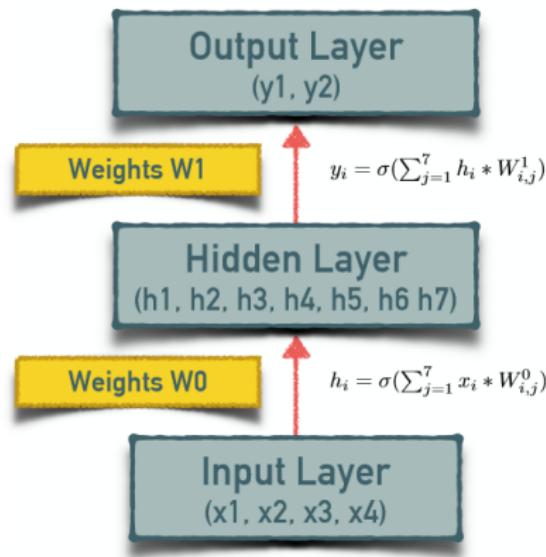
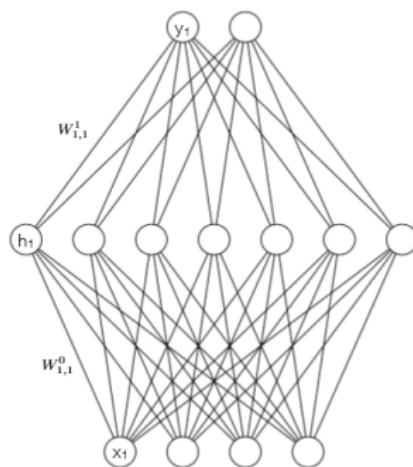
In probability theory, the output of the softmax function can be used to represent a categorical distribution - that is, a probability distribution over K different possible outcomes.

## Softmax function activation

## Multi-Class Classification with NN and SoftMax Function

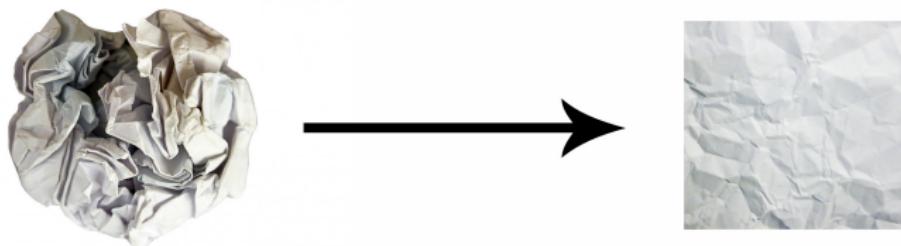


# Feed Forward Neural Network with 1 hidden layer



# Intuition

Complex geometric transformations broken down into simpler ones.



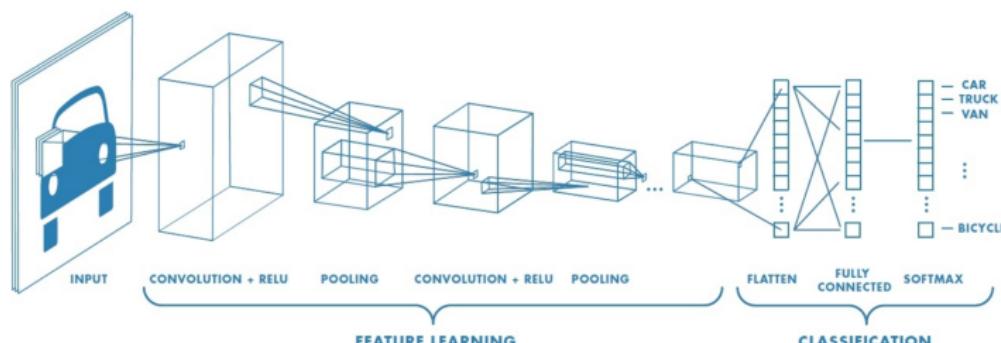
[Tensorflow playground](#)

# CNN

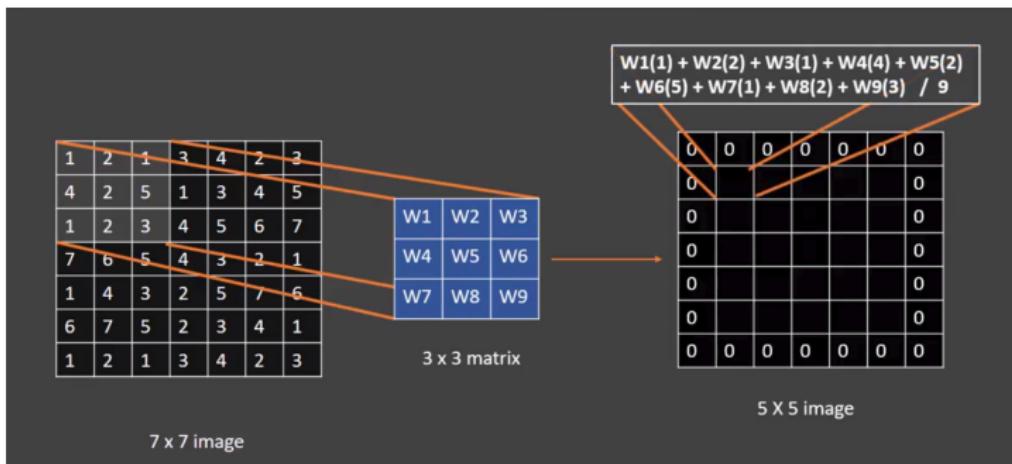
- Similar to Feed Forward Neural Networks.
- Convolutional neural networks (CNN) model can be applied to visual recognition tasks.
- The architecture of a CNN is designed to take advantage of the grid structure of data.
- Hierarchy of representations with increasing level of abstraction.
- Each stage is a kind of trainable feature transform.

# Key components of CNN

- Convolutional layer.
- Pooling.

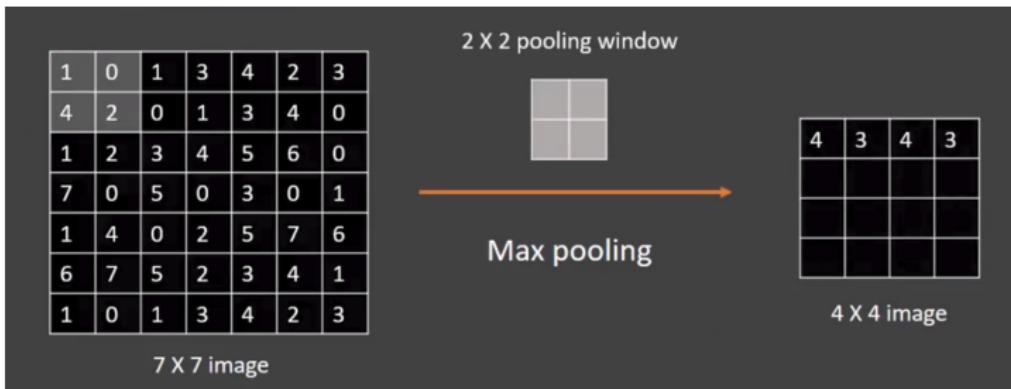


# Convolutional layer



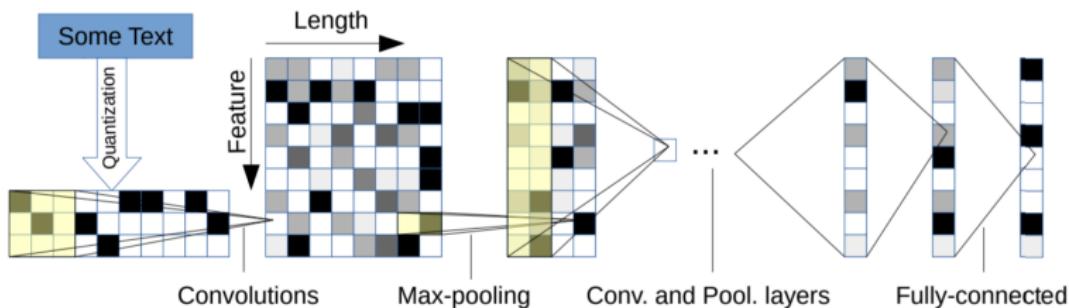
The convolution layer is the main building block of a CNN.

# Pooling layer



Its function is to progressively reduce the spatial size of the representation to reduce the amount of parameters and computation in the network. Pooling layer operates on each feature map independently.

# CNN for text



# RNN

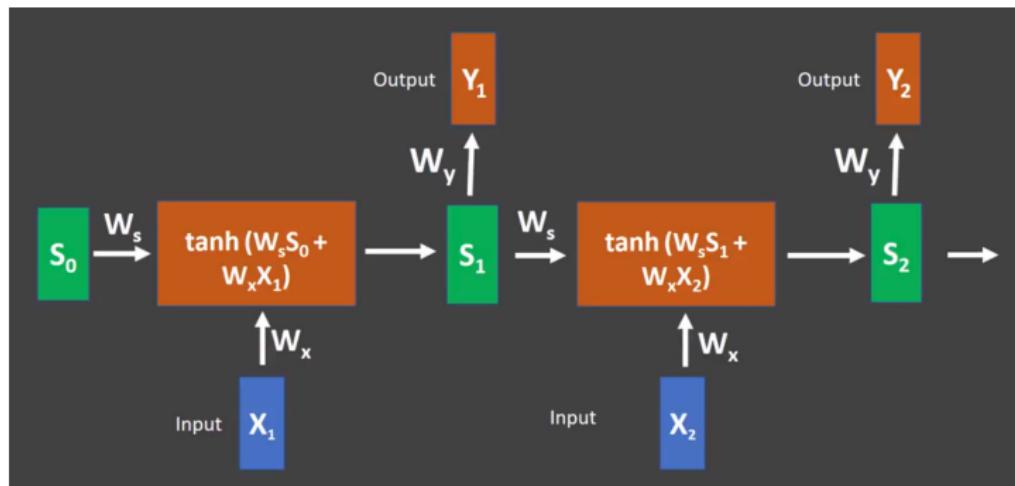
Limitation with Feed Forward Neural Networks (and CNN)...

- Not designed for sequential data.
- Fixed dimension of inputs.
- Does not model memory.

Recurrent neural network...

- To map predictions in sequences.
- This creates an internal state of the network which allows it to exhibit dynamic temporal relations.
- Applications: Image captioning, Sentiment analysis, Question answering, Speech recognition, Time series, Generating music, Machine translation, Self-driving cars, ...

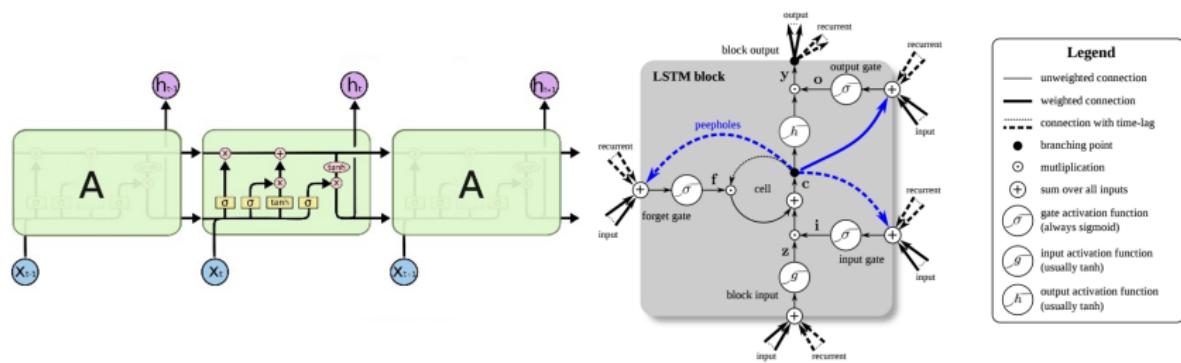
# Vanilla RNN (Unrolled)



Vanilla RNN is affected by the Vanishing Gradient problem and the solution to the problem which is Long Short Term Memory networks (LSTM).

# LSTM

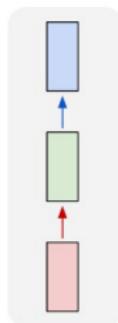
LSTMs are explicitly designed to avoid the long-term dependency problem...



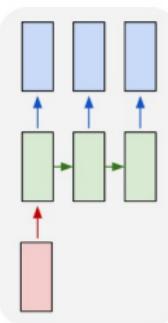
*"Google will soon be a big LSTM"*, Jürgen Schmidhuber

# Input/output size

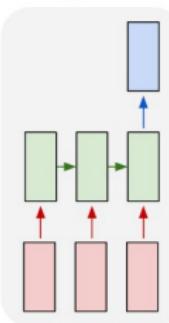
one to one



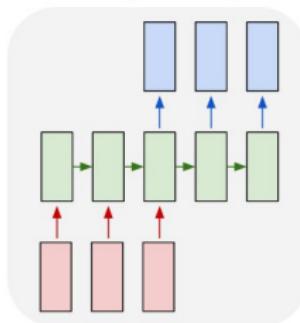
one to many



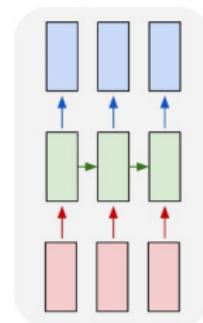
many to one



many to many



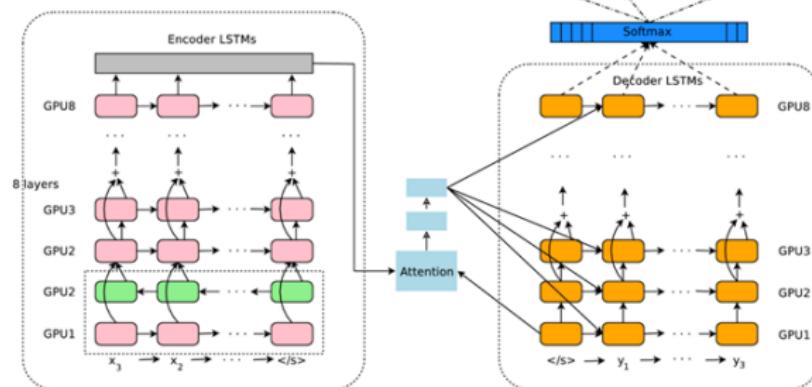
many to many



Operate over sequences of vectors: Sequences in the input, the output, or in the most general case both.

# LSTM with multiple hidden layers

Seq2seq learning for NMT...



# Computer requirements

## MODEL COMPLEXITY IS EXPLODING

7 ExaFLOPS  
60 Million Parameters

20 ExaFLOPS  
300 Million Parameters

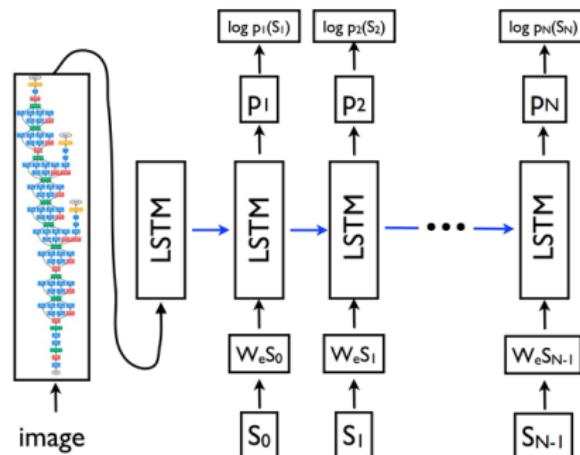
105 ExaFLOPS  
8.7 Billion Parameters

2015 – Microsoft ResNet

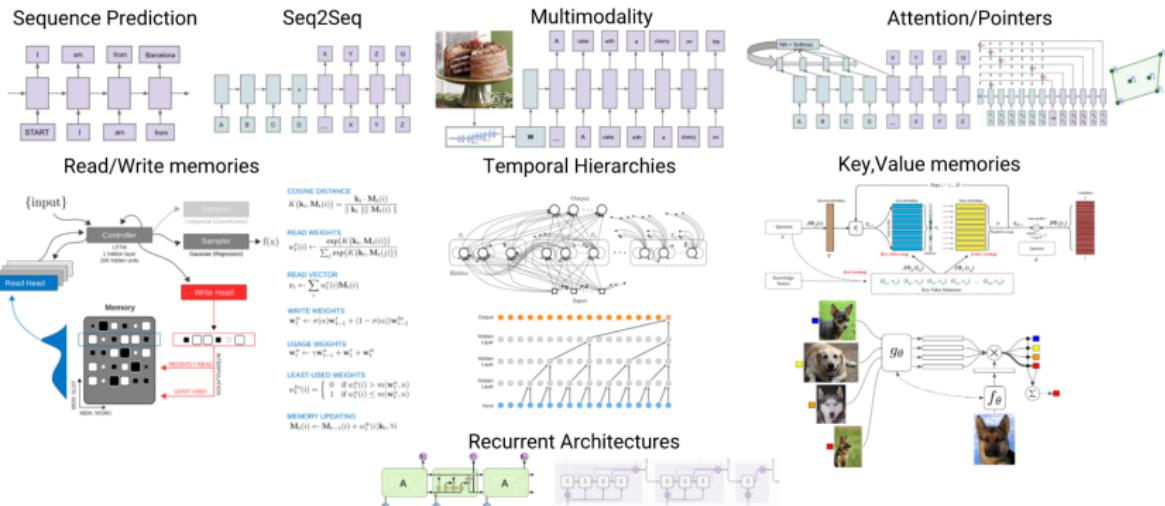
2016 – Baidu Deep Speech 2

2017 – Google NMT

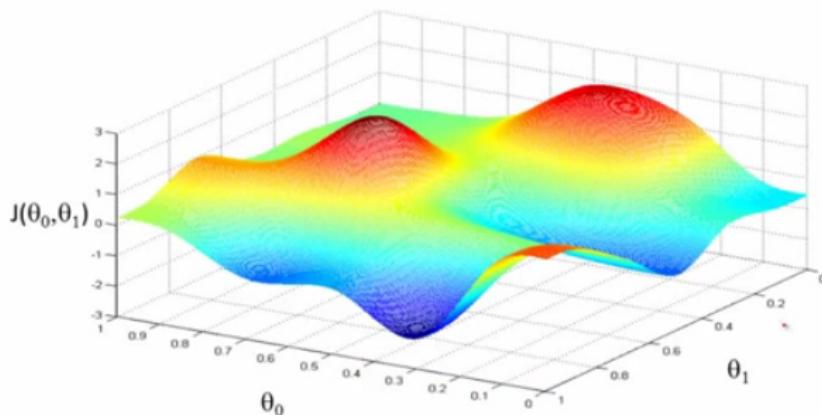
# Image Captioning



# Architectures of Deep Networks



# Loss functions in Deep Learning



Non-convex optimization problem

# Gradient descent

- Gradient descent is a first-order optimization algorithm.
- Find a local minimum of a function using gradient descent, one takes steps proportional to the negative of the gradient of the function at the current point.

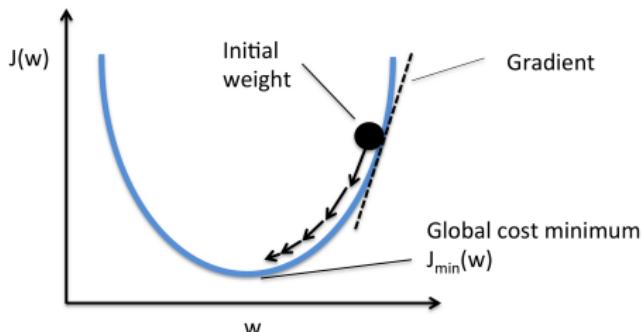


Illustration of gradient descent on a series of level sets.

## Idea to train NN

- The loss function takes the predictions of the network and the true targets and computes a distance score, capturing how well the network has done.
- The fundamental trick in deep learning is to use this score as a feedback signal to adjust the value of the weights a little, in a direction that will lower the loss score for the current batch of examples (**negative gradient**).
- The gradient ( $\nabla J(\theta)$ ) are calculated by **backpropagation**. Calculates the gradient of a loss function with respect to all the weights in the network (**chain rule**) and uses it to update the weights to minimize the loss function.

## Stochastic gradient descent (SGD)

Mini-batch gradient descent performs an update for every mini-batch of  $n$  training examples:

$$\theta_{t+1} = \theta_t - \eta \nabla J(\theta_t / x^{(i:i+n)}; y^{(i:i+n)})$$

Where  $\eta$  is the learning rate.

Can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries that make computing the gradient a mini-batch very efficient. Common mini-batch sizes range between 32 and 256, but can vary for different applications.

**Epoch** is defined as one complete iteration. Fast with GPU computation.

Variants: RMSprop, Adam, ...

# Learning rate

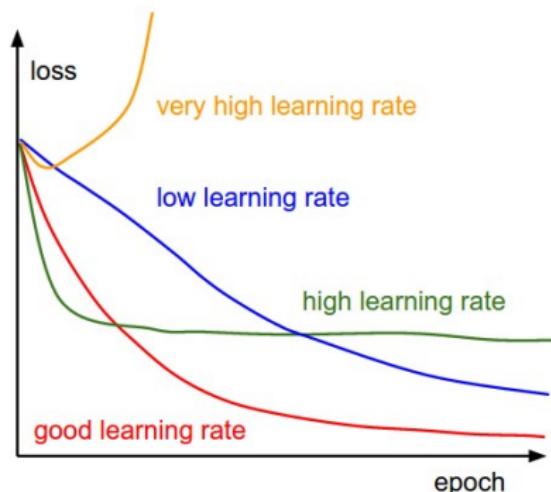


Illustration of gradient descent on a series of level sets.

# Random Initialization Methods

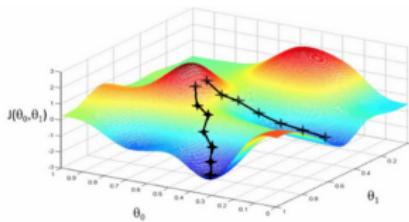
- Initialization with random variables with expectation 0 and low variance.

Xavier's initialization is based on two hypotheses:

- The variance is constant in both the forward and backward directions.
- A dense neural network with symmetric activation  $f$  will have  $f'(y_l) = 1$  when  $y_l = 0$ .

$$\text{Then, } \text{Var} [\theta_l] \simeq \frac{\sqrt{6}}{n_l + n_{l+1}}$$

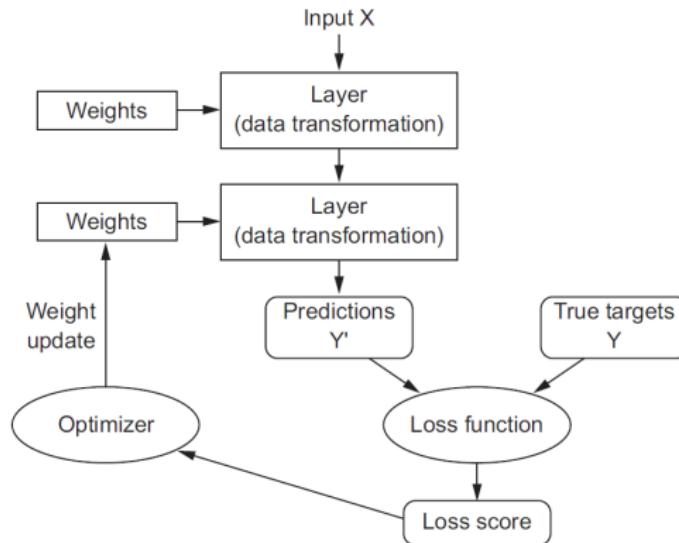
# The Loss Surfaces of Multilayer Networks



*"Neural nets earned a reputation of being finicky and unreliable, which in part caused the community to focus on simpler method with convex loss functions, such as kernel machines and boosting."*

*"...while local minima are numerous, they are relatively easy to find, and they are all more or less equivalent in terms of performance on the test set..."*

# Summary



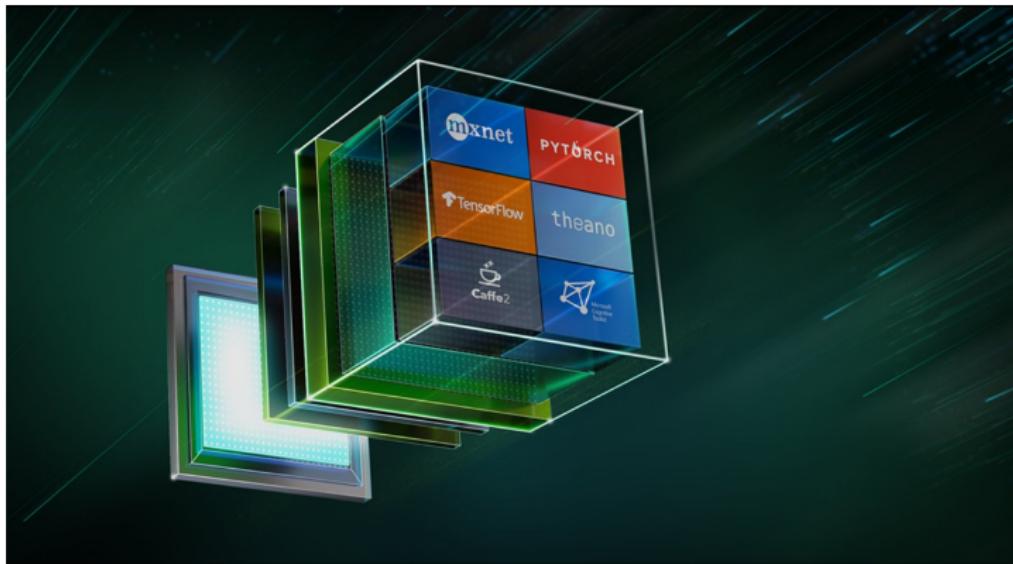
# Some problems with deep learning models

Some limitations...

- Not easily interpreted (not a black box!!).
- Can be brittle (adversarial examples).
- Typically require large amounts of data to perform well.
- Are often very computationally expensive and difficult to train.

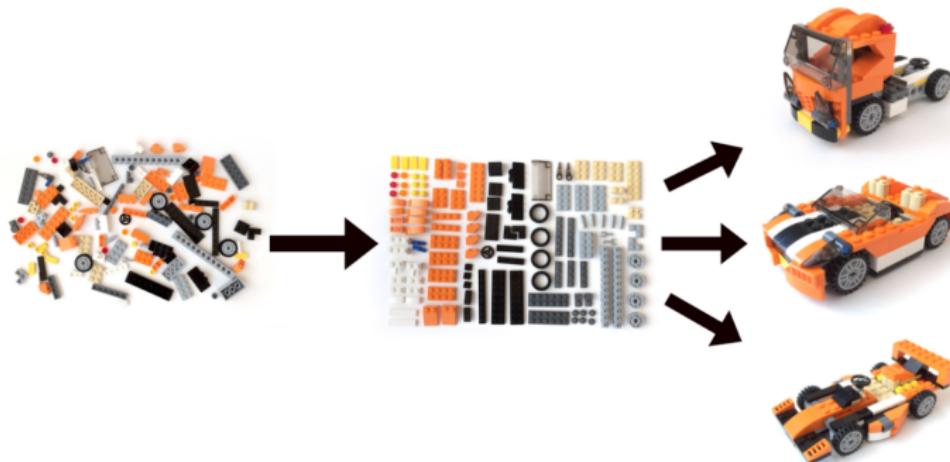
# TOOLS & EXAMPLES

# Backends for deep learning



# Backends for deep learning

## Deep Learning Building Blocks



# TensorFlow: An open-source machine learning framework for everyone



The screenshot shows the official TensorFlow website. The top navigation bar includes links for Install, Develop, API v1.8, Deploy, Extend, Community, Versions, Ecosystem, a search bar labeled 'Buscar', and GitHub. Below the navigation is a large orange banner with the text 'An open source machine learning framework for everyone' and a 'GET STARTED' button.



TensorFlow 1.8 has arrived!

We're excited to announce the release of TensorFlow 1.8! Check out the announcement to upgrade your code with ease.

[LEARN MORE](#)



TensorFlow Dev Summit 2018

Thousands of people from the TensorFlow community participated in the second TensorFlow Dev Summit. Watch the keynote and talks now.

[WATCH NOW](#)



Announcing TensorFlow.js!

Learn more about our new library for machine learning in the browser using JavaScript.

[LEARN MORE](#)

<https://www.tensorflow.org/>

# TensorFlow: An open-source machine learning framework for everyone

The screenshot shows the TensorFlow website's 'Companies using TensorFlow' section. The page has a header with navigation links like 'Install', 'Develop', 'API v1.8', 'Deploy', 'Extend', 'Community', 'Versions', 'Ecosystems', a search bar, and a GitHub link. Below the header is a grid of company logos:

Row 1	airbnb	AMD	NVIDIA	UBER
Row 2	SAP	kakao	DeepMind	Dropbox
Row 3	ebay	Google	Snaps	intel
Row 4	Coca-Cola	MI	QUALCOMM	Twitter
Row 5	INSGO	SDR	AIRBUS DEFENCE & SPACE	ARM
Row 6	caicloud	CastBox	CEVA	cheetahmobile
Row 7	CIST	IBM Power Systems	JD.COM 京东	Lenovo

# TensorFlow: An open-source machine learning framework for everyone

The screenshot shows the GitHub repository page for TensorFlow. At the top, it displays basic repository statistics: 30,843 commits, 26 branches, 53 releases, 1,409 contributors, and an Apache-2.0 license. Below this, a list of recent commits is shown, each with a user icon, author, commit message, and timestamp. The commits are from various contributors like imsherdan, sb2nov, tensorflow, third\_party, tools, util/python, gitignore, ACKNOWLEDGMENTS, ADOPTERS.md, AUTHORS, BUILD, and CODEOWNERS. The timestamps range from 17 hours ago to 2 years ago.

Author	Commit Message	Timestamp
imsherdan and sb2nov	Support a single Tensor in StagingArea.put()	17 hours ago
tensorflow	Support a single Tensor in StagingArea.put()	17 hours ago
third_party	Resolve conflicts.	3 days ago
tools	Allow to download clang and use clang for CPU builds.	10 days ago
util/python	Merge changes from github.	4 months ago
gitignore	Automated g4 rollback of changelist 179260538	4 months ago
ACKNOWLEDGMENTS	TensorFlow: Improve performance of Alexnet	2 years ago
ADOPTERS.md	Internal file cleanup.	2 years ago
AUTHORS	Merge changes from github.	4 months ago
BUILD	Make LICENSE visible to bazel.	3 months ago
CODEOWNERS	Merge changes from github.	3 months ago

<https://github.com/tensorflow/tensorflow>

# TensorFlow/Keras in GitHub

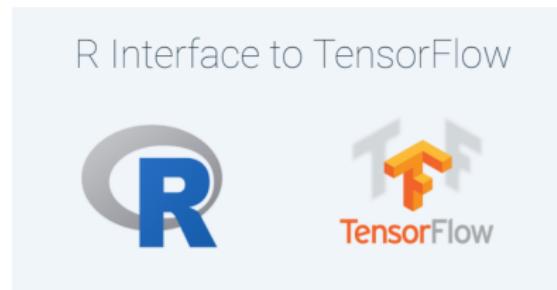
Top libraries by GitHub issues opened

#1:	8370	tensorflow/tensorflow
#2:	5806	fchollet/keras
#3:	4558	dmlc/mxnet
#4:	3908	BVLC/caffe
#5:	2465	Theano/Theano
#6:	2462	baidu/paddle
#7:	2264	deeplearning4j/deeplearning4j
#8:	2124	Microsoft/CNTK
#9:	1601	pytorch/pytorch
#10:	1139	NVIDIA/DIGITS
#11:	1005	pfnet/chainer
#12:	738	caffe2/caffe2
#13:	709	tflearn/tflearn
#14:	664	davisking/dlib
#15:	575	torch/torch7
#16:	488	Lasagne/Lasagne
#17:	469	clab/dynet
#18:	324	NervanaSystems/neon
#19:	47	deepmind/sonnet

Top libraries by GitHub stars

#1:	71627	tensorflow/tensorflow
#2:	20489	BVLC/caffe
#3:	20038	fchollet/keras
#4:	12558	Microsoft/CNTK
#5:	11369	dmlc/mxnet
#6:	7712	pytorch/pytorch
#7:	7332	torch/torch7
#8:	7297	deeplearning4j/deeplearning4j
#9:	6981	Theano/Theano
#10:	6767	tflearn/tflearn
#11:	5742	caffe2/caffe2
#12:	5544	baidu/paddle
#13:	5336	deeplearning/sonnet
#14:	3242	Lasagne/Lasagne
#15:	3232	NervanaSystems/neon
#16:	2987	pfnet/chainer
#17:	2833	davisking/dlib
#18:	2525	NVIDIA/DIGITS
#19:	1775	clab/dynet

# TensorFlow



<https://tensorflow.rstudio.com/>

## Install from github

```
> library("devtools")
> devtools::install_github("rstudio/tensorflow")
> devtools::install_github("rstudio/keras")
```

## Tensorflow example

### tfSession()

```
> library(tensorflow)
> sess <- tf$Session()
> hello <- tf$constant('Hello, TensorFlow!')
> sess$run(hello)

b'Hello, TensorFlow!'
```

### Custom linear regression example

```
> # Create 100 phony x, y data points, y = x * 0.1 + 0.3
> x_data <- runif(100, min=0, max=1)
> y_data <- x_data * 0.1 + 0.3
> # Try to find values for W and b that compute y_data = W * x_data + b
> # (We know that W should be 0.1 and b 0.3, but TensorFlow will
> # figure that out for us.)
> W <- tf$Variable(tf$random_uniform(shape(1L), -1.0, 1.0))
> b <- tf$Variable(tf$zeros(shape(1L)))
> y <- W * x_data + b
> # Minimize the mean squared errors.
> loss <- tf$reduce_mean((y - y_data) ^ 2)
> optimizer <- tf$train$GradientDescentOptimizer(0.5)
> train <- optimizer$minimize(loss)
```



# Tensorflow example

## tfSession()

```
> # Launch the graph and initialize the variables.  
> sess = tf$Session()  
> sess$run(tf$global_variables_initializer())  
> # Fit the line (Learns best fit is W: 0.1, b: 0.3)  
> for (step in 1:201) {  
+   sess$run(train)  
+   if (step %% 20 == 0)  
+     cat(step, "-", sess$run(W), sess$run(b), "\n")  
+ }  
  
20 - -0.01607878 0.3667336  
40 - 0.07722016 0.3130961  
60 - 0.09552959 0.30257  
80 - 0.09912271 0.3005044  
100 - 0.09982784 0.300099  
120 - 0.0999662 0.3000194  
140 - 0.09999336 0.3000038  
160 - 0.0999987 0.3000008  
180 - 0.09999975 0.3000001  
200 - 0.09999989 0.3000001
```

# Keras: User-friendly API which makes it easy to quickly prototype deep learning models

<https://keras.rstudio.com/>



---

R interface to Keras

It contains about  $\simeq 500$  functions (v.2.1.6)

# Imperative API

## k-eval

```
> library(keras)
> x <- c(1:15)
> prod_ac <- k_cumprod(x, axis = 1)
> k_eval(prod_ac)

[1]      1      2      6     24     120      720
[7]  5040   40320  362880  3628800 39916800 479001600
[13] 1932053504 1278945280 2004310016

> sig <- k_sigmoid(c(-3, -2, -1, 0, 1, 2, 3))
> k_eval(sig)

[1] 0.04742587 0.11920292 0.26894143 0.50000000 0.73105860 0.88079703 0.95257413

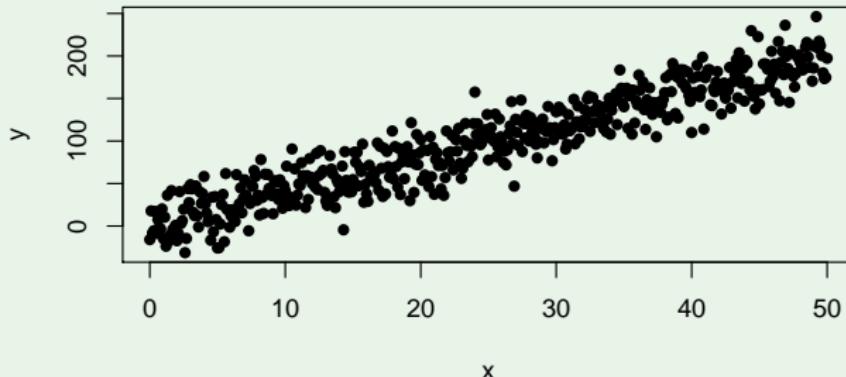
> k_eval(k_random_normal(c(4, 6), mean = 0, stddev = 1))

      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
[1,] -1.212058  0.08014527 -1.3834615  0.9868450 -0.93238080  0.94287306
[2,]  1.536298 -0.68395638 -0.3206587 -0.2576625 -1.44378328  1.15996063
[3,]  1.337245 -1.34764791 -0.4571670 -0.6830830  0.01222763 -0.07862066
[4,] -0.198133  1.04633021 -0.6374950 -0.2465717 -0.35541692  0.33105391
```

# Data

## simulated data

```
> x <- seq(0, 50, by = 0.1)
> y <- 4*x + rnorm(n = length(x), 0, 20)
> plot(x, y, pch=16)
```



# Model definition: Linear Regression

- In Keras, we use its own data type `keras-model-sequential()` to configure the network.
- Without hidden layers configuration.
- For the output, we set the activation to linear regression to get the value prediction.

## keras-model-sequential

```
> linear_regression <- keras_model_sequential() %>%
+   layer_dense(units = 1, activation = 'linear', input_shape = c(1))
```

# Model training

## summary

```
> summary(linear_regression)
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 1)	2

Total params: 2  
Trainable params: 2  
Non-trainable params: 0

And the parameters learned...

## Model compile and training

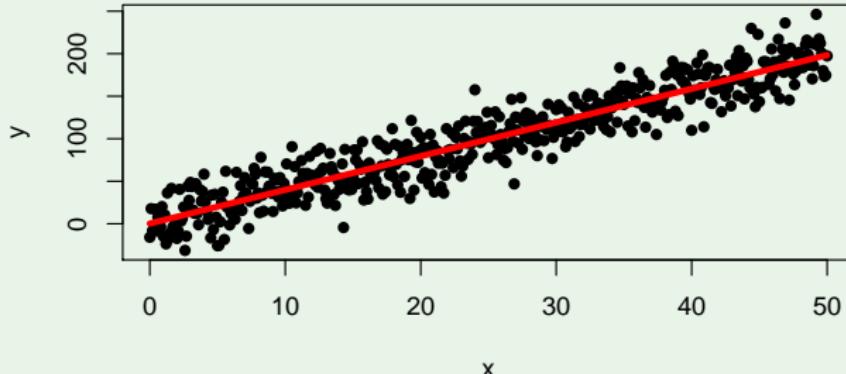
```
> linear_regression %>% compile(loss = 'mean_squared_error', optimizer = optimizer_rmsprop())  
> # Fit model to data  
> history <- linear_regression %>% fit(x, y, batch_size = 256, epochs = 5000, verbose = 0)  
> unlist(get_weights(linear_regression))  
  
[1] 3.9628351 0.2595091
```



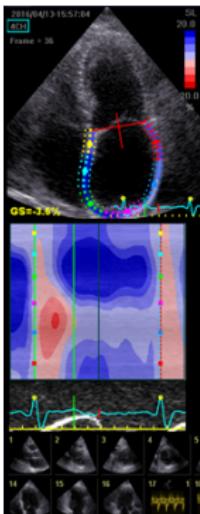
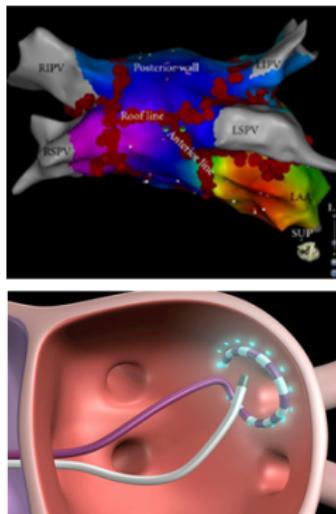
# Prediction

## predict

```
> #Model predictions
> prediction <- linear_regression %>% predict(x)
> plot(x, y, pch=16)
> points(x, prediction, type='l', col='red', lwd=4)
```



# Paper STRAIN



Atrial Fibrillation ablation and Strain variables

# Model architecture: Logistic Regression

- Without hidden layers configuration.
- For the output, we set the activation to softmax to get a probabilistic prediction for each category.

## layer-dense

```
> logistic_regression <- keras_model_sequential() %>%  
+   layer_dense(units = 1, activation = 'sigmoid', input_shape = c(21))
```

# Model training

We train the model over 2500 epochs.

## Model prediction

```
> logistic_regression %>% compile(loss = 'binary_crossentropy', optimizer = optimizer_adam())
> # Fit model to data
> history <- logistic_regression %>% fit(x, y, batch_size = 64, epochs = 2500, verbose = 0)
> score <- logistic_regression %>% evaluate(x, y, verbose = 0)
> score

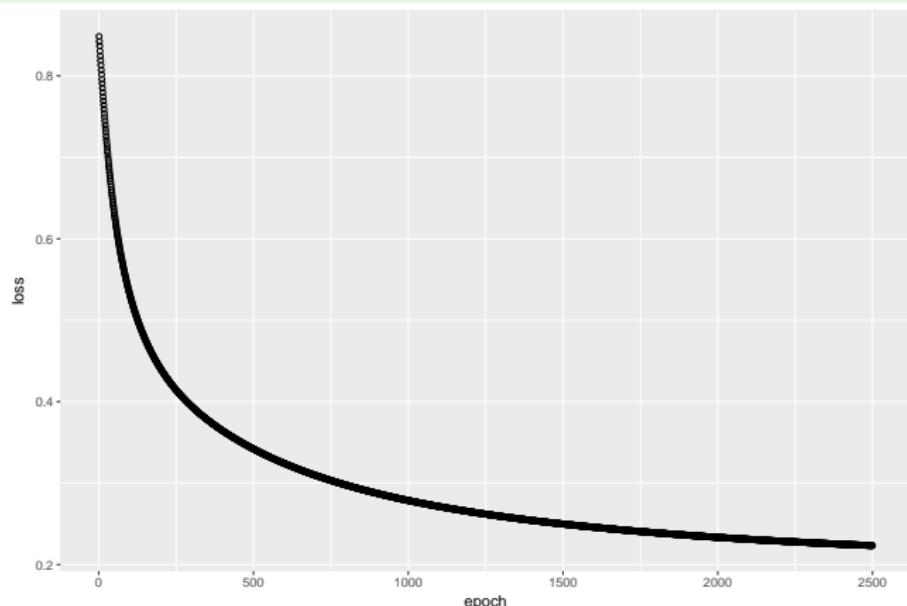
      loss
0.2233094

> class_prediction <- logistic_regression %>% predict_classes(x)
> confusion_matrix <- table(class_prediction, y)
> sum(diag(confusion_matrix))/sum(confusion_matrix)*100

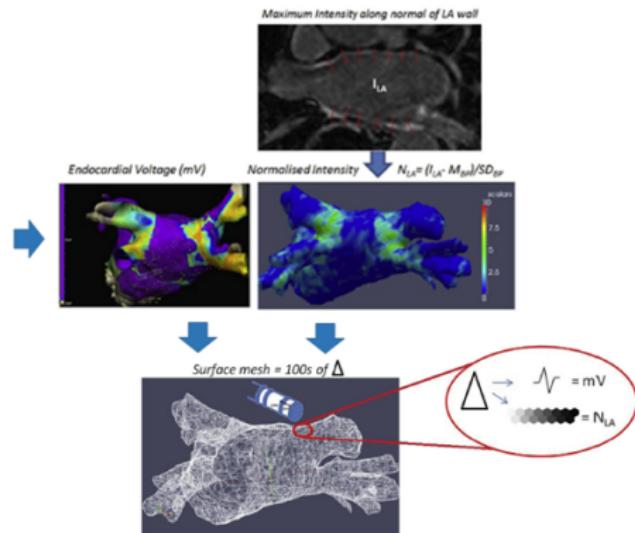
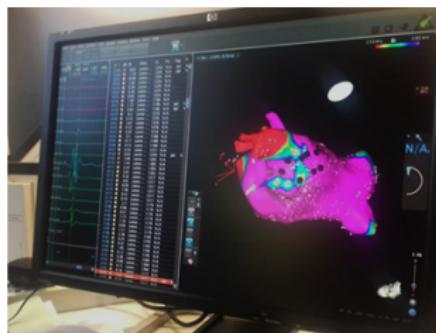
[1] 91.15044
```

# Monitoring the error

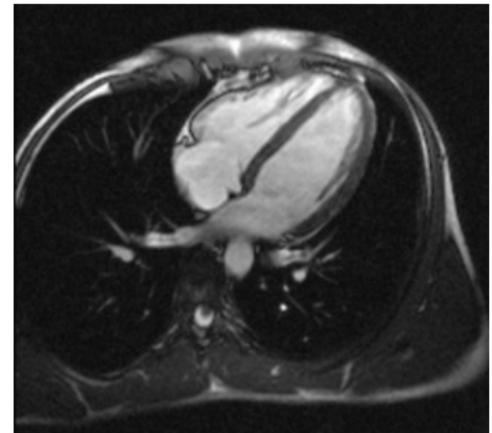
## Training curve



# RMN to IIR prediction

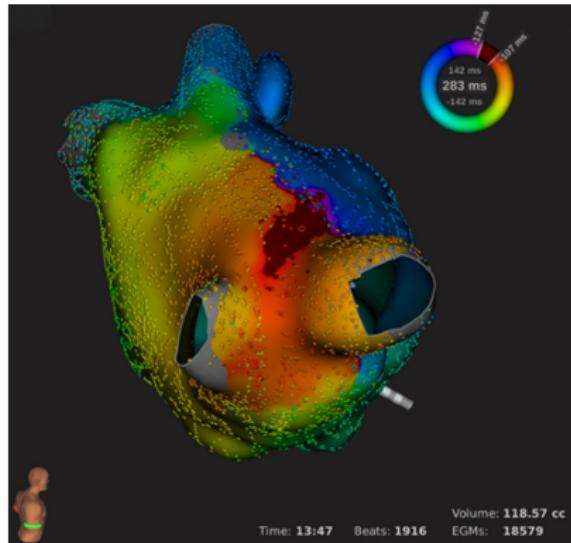


# RMN to IIR prediction



Cardiac magnetic resonance image

## RMN to IIR prediction



An electroanatomical mapping system is a useful tool for complex arrhythmia interventions.

## RMN to IIR prediction

- Model exploration is mandatory...
- Linear models not always the best choice, voltage negative predictions for this problem!
- We need hidden layers configuration.

### NNET model architecture

```
> nnet_rmn_iir <- keras_model_sequential()
> nnet_rmn_iir %>%
+   layer_dense(units = 128, input_shape = c(27)) %>%
+   layer_activation('relu') %>%
+   layer_dense(units = 64) %>%
+   layer_activation('relu') %>%
+   layer_dense(units = 1)
```

# Computation graph

## summary

```
> summary(nnet_rmn_iir)
```

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 128)	3584
activation_1 (Activation)	(None, 128)	0
dense_4 (Dense)	(None, 64)	8256
activation_2 (Activation)	(None, 64)	0
dense_5 (Dense)	(None, 1)	65

Total params: 11,905

Trainable params: 11,905

Non-trainable params: 0

After 10,000 iterations with a few thousand parameters, the model gets 0.72 mv of *MAE* in the test set.

## Real-world examples of data tensors

- *Vector data*—2D tensors of shape (samples, features)
- *Timeseries data or sequence data*—3D tensors of shape (samples, timesteps, features)
- *Images*—4D tensors of shape (samples, height, width, channels) or (samples, channels, height, width)
- *Video*—5D tensors of shape (samples, frames, height, width, channels) or (samples, frames, channels, height, width)

# Kaggle competition



## State Farm Distracted Driver Detection

Tue 5 Apr 2016

Merger and 1st Submission Deadline

Mon 1 Aug 2016 (38 days to go)

Dashboard

Home Data Make a submission

Information

Description Evaluation Rules Prizes Timeline

Forum

Scripts

New Script New Notebook

Leaderboard

My Team

My Submissions

Competition Details » Get the Data » Make a submission

### Can computer vision spot distracted drivers?

We've all been there: a light turns green and the car in front of you doesn't budge. Or, a previously unremarkable vehicle suddenly slows and starts swerving from side-to-side.

When you pass the offending driver, what do you expect to see? You certainly aren't surprised when you spot a driver who is texting, seemingly enraptured by social media, or in a lively hand-held conversation on their phone.



# Image example

## readImage

```
> library("EBImage")
> image_c0 <- readImage("c0_img_2093.jpg")
> display(image_c0, method = "raster")
```



# Image example

readImage

```
> image_c2 <- readImage("c2_img_13859.jpg")
> display(image_c2, method = "raster")
```



# Data

## data.matrix

```
> Driver_Detection_sample[1:4, 1:7]

 target      X1      X2      X3      X4      X5      X6
 1     3 0.1200282 0.1474765 0.1778427 0.1883918 0.1921196 0.2078043
 2     0 0.1572761 0.1533549 0.1611973 0.1611973 0.1729608 0.1729608
 3     1 0.1500741 0.1293384 0.1293384 0.1214961 0.1336643 0.1532702
 4     6 0.1949529 0.1650337 0.1689125 0.1598137 0.1650761 0.1716616

> dim(Driver_Detection_sample)

[1] 5000 1025

> sample_index <- sample(1:nrow(Driver_Detection_sample), size=nrow(Driver_Detection_sample)*0.8)
> train <- data.matrix(Driver_Detection_sample[sample_index,])
> test <- data.matrix(Driver_Detection_sample[-sample_index,])
> x_train <- train[,-1]; y_train <- train[,1]
> x_test <- test[,-1]; y_test <- test[,1]
```

# Data manipulation with Keras

## array-reshape

```
> img_rows <- 32
> img_cols <- 32
> batch_size <- 128
> num_classes <- 10
> # Redefine dimension of train/test inputs
> x_train <- array_reshape(x_train, c(nrow(x_train), img_rows, img_cols, 1))
> x_test <- array_reshape(x_test, c(nrow(x_test), img_rows, img_cols, 1))
> # Input image dimensions
> input_shape <- c(img_rows, img_cols, 1)
> y_train <- to_categorical(y_train, num_classes)
> y_test <- to_categorical(y_test, num_classes)
```

# CNN definition

**Convolutional network** architecture designed for images.

## CNN model definition

```
> model <- keras_model_sequential()
> model %>% layer_conv_2d(filters = 32, kernel_size = c(3,3),
+   activation = 'relu', input_shape = input_shape) %>%
+   layer_conv_2d(filters = 64, kernel_size = c(3,3), activation = 'relu') %>%
+   layer_max_pooling_2d(pool_size = c(2, 2)) %>%
+   layer_flatten() %>%
+   layer_dense(units = 128, activation = 'relu') %>%
+   layer_dense(units = num_classes, activation = 'softmax')
```

Compile and Train...

## Fit the model

```
> model %>% compile(loss = loss_categorical_crossentropy,
+   optimizer = optimizer_adadelta(), metrics = c('accuracy'))
> model %>% fit(x_train, y_train, batch_size = batch_size,
+   epochs = 15, verbose = 1, validation_data = list(x_test, y_test))
```

# Model parameters

## summary model

```
> summary(model)
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 30, 30, 32)	320
conv2d_2 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 64)	0
flatten_1 (Flatten)	(None, 12544)	0
dense_6 (Dense)	(None, 128)	1605760
dense_7 (Dense)	(None, 10)	1290

Total params: 1,625,866

Trainable params: 1,625,866

Non-trainable params: 0

# Model performance

## predict

```
> scores <- model %>% evaluate(x_test, y_test, verbose = 1)
> scores

$loss
[1] 0.1115962

$acc
[1] 0.969

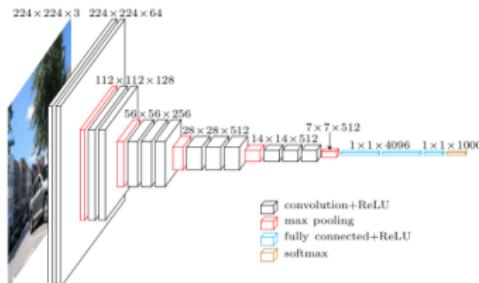
> predicted_class <- model %>% predict_classes(x_test)
> confusion_matrix <- table(predicted_class, test[,1])
> confusion_matrix

predicted_class   0   1   2   3   4   5   6   7   8   9
      0 99  1  0  2  0  0  0  0  1  2
      1  0 97  0  0  0  0  0  0  1  0
      2  2  0 101  1  0  1  3  0  0  0
      3  0  0  0  86  1  0  0  0  1  0
      4  1  0  0  0 109  0  0  2  0  2
      5  0  0  0  0  0  84  0  0  0  0
      6  0  1  1  0  0  1  92  1  0  0
      7  1  0  0  0  0  0  0 108  0  0
      8  0  0  0  0  0  1  0  0  90  1
      9  0  1  0  0  1  0  0  0  1 103
```



# Classify images with a pretrained model

- You can use a pre-trained **VGG** network to predict the class of an image.
- **VGG** network gives the recent state-of-art prediction accuracy on **IMAGENET** dataset.
- Model with 143,667,240 parameters.



## load.image

```
> library(jpeg)
> img_path <- 'D:/SESSIO_DEEP_LEARNING/BEAMER/elephant.jpg'
> ### load the image and display it in R
> jj <- readJPEG(img_path)
> plot(0:1, 0:1, type ="n", ann=FALSE, axes=FALSE)
> rasterImage(jj, 0, 0, 1, 1)
```



# Classify images with a pretrained model

## VGG prediction

```
> img <- image_load(img_path, target_size = c(224, 224))
> x <- image_to_array(img)
> ## ensure we have 4d tensor with single element in the batch dimensions,
> ## the preprocess the input for prediction
> dim(x) <- c(1, dim(x))
> x <- imagenet_preprocess_input(x)
> modelvgg19 <- application_vgg19(weights = 'imagenet')
> pred_vgg19 <- modelvgg19 %>% predict(x)
> ppvgg <- imagenet_decode_predictions(pred_vgg19, top =10)[[1]]
> ppvgg[1:5,]

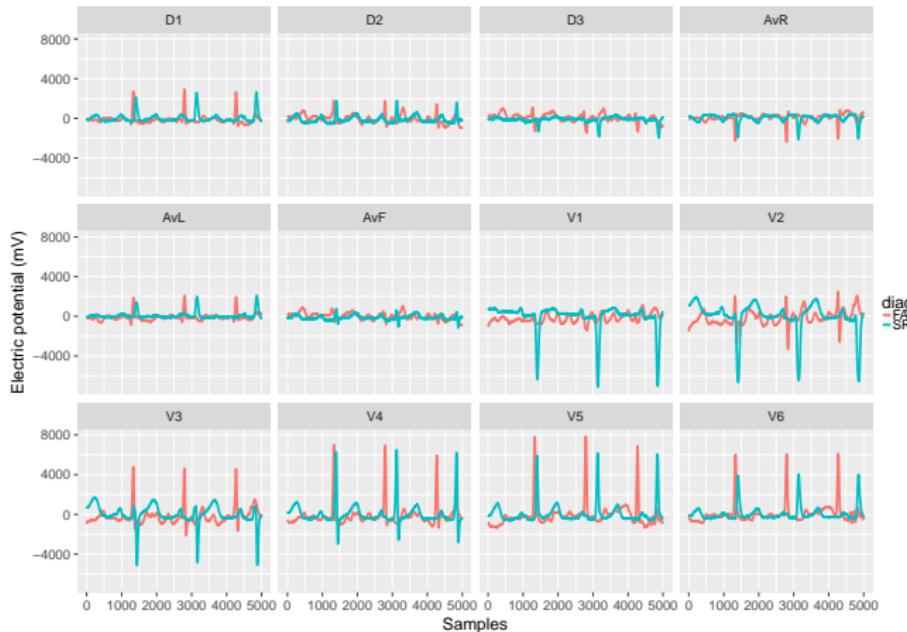
  class_name class_description      score
1 n01871265          tusker 0.4945689142
2 n02504458  African_elephant 0.4688400328
3 n02504013 Indian_elephant 0.0330304392
4 n02412080            ram 0.0007902890
5 n02454379   armadillo 0.0006282972
```

It's a tusker!! :)

# ECG monitoring example



# ECG monitoring example



# ECG monitoring example

## layer-lstm

```
> model_ecg <- keras_model_sequential() %>%
+   layer_lstm(units = 32, return_sequences = TRUE, input_shape = c(1, 12)) %>%
+   layer_lstm(units = 32, return_sequences = TRUE) %>%
+   layer_lstm(units = 32) %>%
+   layer_dense(units = 1, activation = "sigmoid")
```

Accuracy 94% in the validation set.

# IMDB database



Dataset of 25,000 movies reviews from IMDB, labeled by sentiment (positive/negative).

# IMDB database

## Data to dictionary

```
> max_features <- 2000
> max_len <- 500
> imdb <- dataset_imdb(num_words = max_features)
> imdb$train$x[[10]]
```

[1]	1	14	20	47	111	439	2	19	12	15	166	12	216	125	40
[16]	6	364	352	707	1187	39	294	11	22	396	13	28	8	202	12
[31]	1109	23	94	2	151	111	211	469	4	20	13	258	546	1104	2
[46]	12	16	38	78	33	211	15	12	16	2	63	93	12	6	253
[61]	106	10	10	48	335	267	18	6	364	1242	1179	20	19	6	1009
[76]	7	1987	189	5	6	2	7	2	2	95	1719	6	2	7	2
[91]	2	49	369	120	5	28	49	253	10	10	13	1041	19	85	795
[106]	15	4	481	9	55	78	807	9	375	8	1167	8	794	76	7
[121]	4	58	5	4	816	9	243	7	43	50					

## Pad sequences

```
> c(c(x_train, y_train), c(x_test, y_test)) %<-% imdb
> x_train <- pad_sequences(x_train, maxlen = max_len)
> x_test <- pad_sequences(x_test, maxlen = max_len)
```

# IMDB database

## Model

```
> model <- keras_model_sequential() %>%
+   layer_embedding(input_dim = max_features, output_dim = 128,
+                   input_length = max_len, name = "embed") %>%
+   layer_conv_1d(filters = 32, kernel_size = 7, activation = "relu") %>%
+   layer_max_pooling_1d(pool_size = 5) %>%
+   layer_conv_1d(filters = 32, kernel_size = 7, activation = "relu") %>%
+   layer_global_max_pooling_1d() %>%
+   layer_dense(units = 1, activation = "sigmoid")
> model %>% compile(optimizer = "rmsprop", loss = "binary_crossentropy", metrics = c("acc"))
```

Accuracy 86% in the test set.

# Advanced features

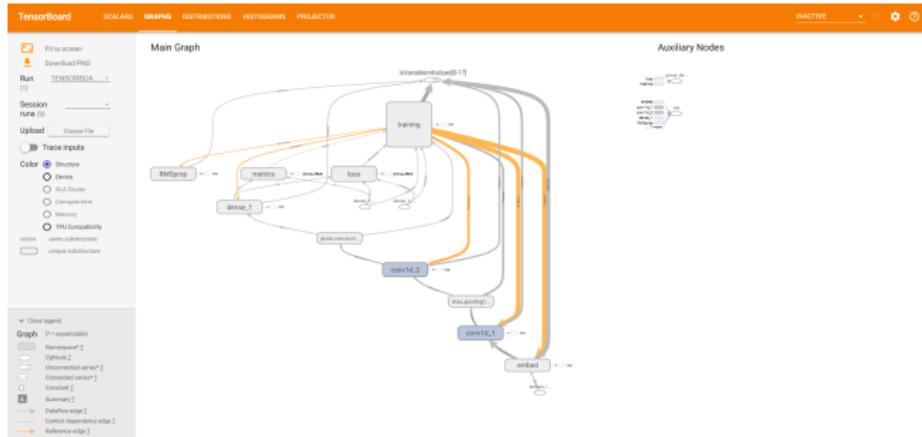
## Advanced features in Keras

- Tensorboard visualization
- Deep Art
- Multitasking learning
- Model to production and google cloud integration
- Generative Adversarial Learning, Reinforcement learning (alpha Go), Transfer Learning k
- And many others to have fun!... :)

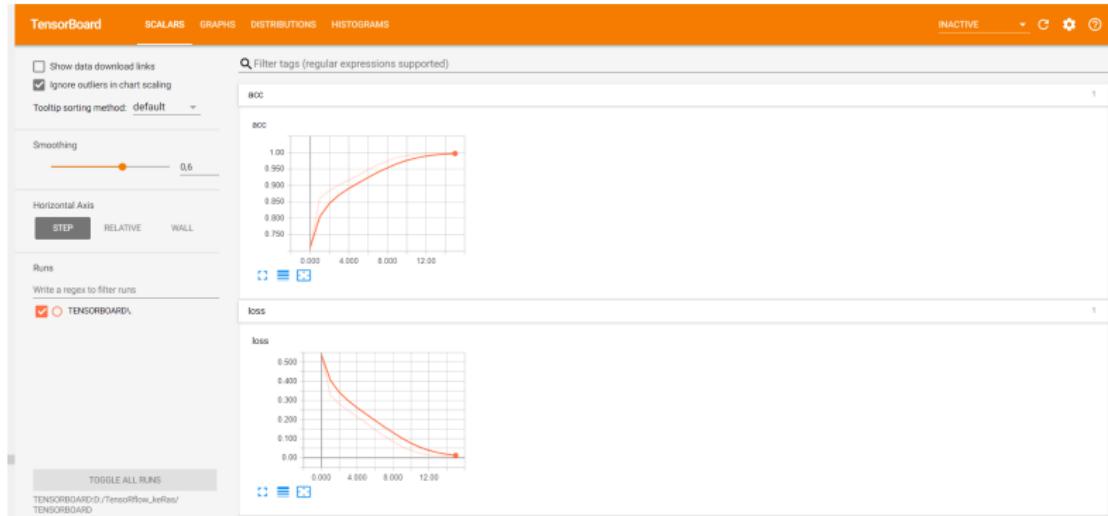
# TENSORBOARD visualization

## Model visualization

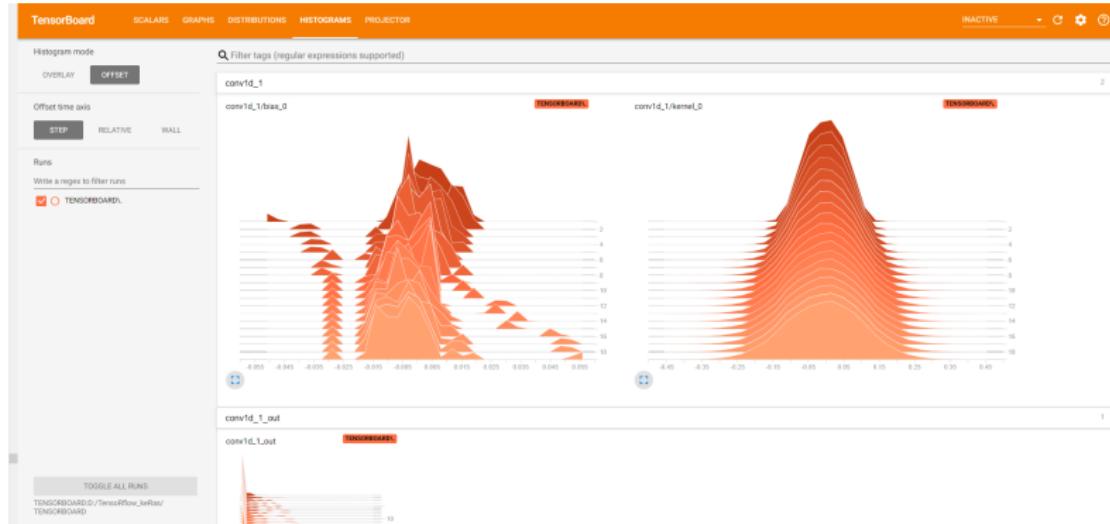
```
> callbacks <- list(callback_tensorboard(log_dir =  
+                               "D:/TensoRflow_keRas/TENSORBOARD", histogram_freq = 1, embeddings_freq = 1))  
> history <- model %>% fit(x_train, y_train, epochs = 10, batch_size = 128,  
+                               validation_split = 0.2, callbacks = callbacks)  
> tensorboard("D:/TensoRflow_keRas/TENSORBOARD")
```



# TENSORBOARD visualization



# TENSORBOARD visualization



# TENSORBOARD visualization



# Neural Style Transfer



Neural Style Transfer in R

# Neural Style Transfer

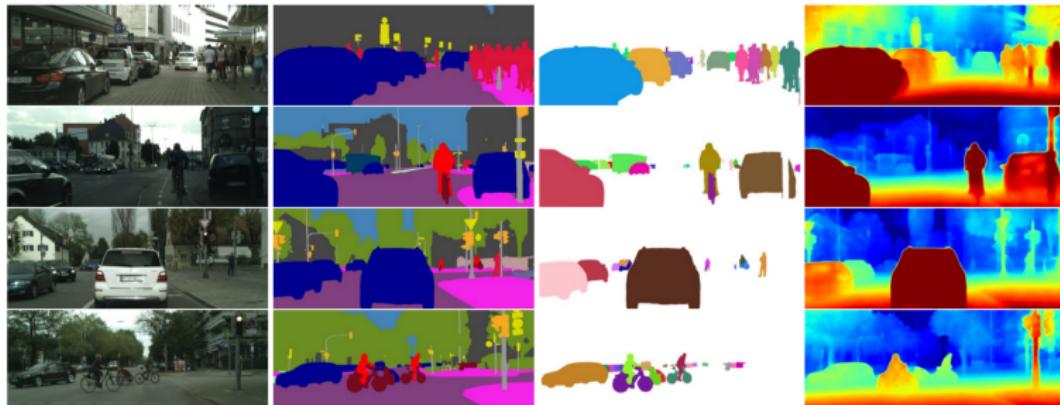
## Model iteration

```
> for (i in 1:iterations) {  
+  
+   # Runs L-BFGS over the pixels of the generated image to minimize the neural style loss.  
+   opt <- optim(  
+     array_reshape(x, dim = length(x)),  
+     fn = evaluator$loss,  
+     gr = evaluator$grads,  
+     method = "L-BFGS-B",  
+     control = list(maxit = 15)  
+   )  
+  
+   cat("Loss:", opt$value, "\n")  
+  
+   image <- x <- opt$par  
+   image <- array_reshape(image, dms)  
+  
+   im <- deprocess_image(image)  
+   plot(as.raster(im))  
+ }
```

# Neural Style Transfer



# Multi-task learning



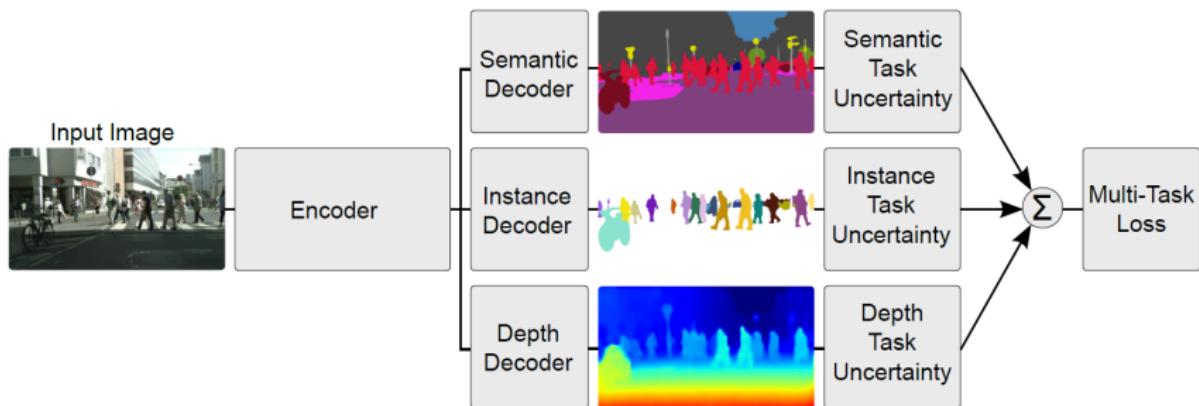
(a) Input image

(b) Segmentation output

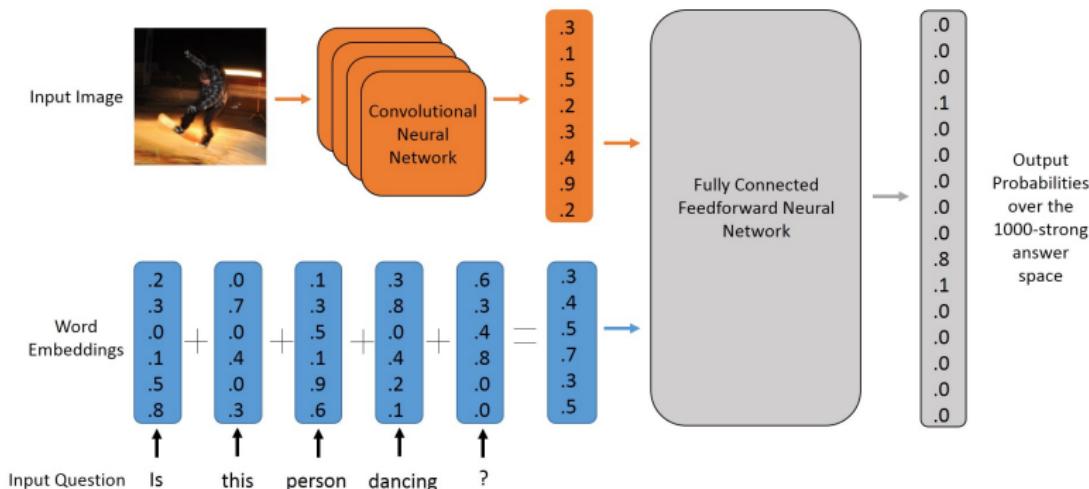
(c) Instance output

(d) Depth output

# Multi-task learning



## Multi-task learning



# Multi-task learning

## Functional API implementation of a two-input question-answering model

Embeds the inputs into a sequence of vectors of size 64

The text input is a variable-length sequence of integers. Note that you can optionally name the inputs.

Same process (with different layer instances) for the question

```
library(keras)

text_vocabulary_size <- 10000
ques_vocabulary_size <- 10000
answer_vocabulary_size <- 500

text_input <- layer_input(shape = list(NULL),
                           dtype = "int32", name = "text")

encoded_text <- text_input %>%
  layer_embedding(input_dim = 64, output_dim = text_vocabulary_size) %>%
  layer_lstm(units = 32)

question_input <- layer_input(shape = list(NULL),
                               dtype = "int32", name = "question")

encoded_question <- question_input %>%
  layer_embedding(input_dim = 32, output_dim = ques_vocabulary_size) %>%
  layer_lstm(units = 16)

concatenated <- layer_concatenate(list(encoded_text, encoded_question))
```

# Other libraries in the ecosystem

## Install from github

```
> devtools::install_github("rstudio/cloudml")
> devtools::install_github("rstudio/tfdeploy")
> devtools::install_github("rstudio/tfestimators")
```

## Train

Train models with [tensorflow.rstudio.com](https://tensorflow.rstudio.com)



# CONCLUSIONS

## Final thought

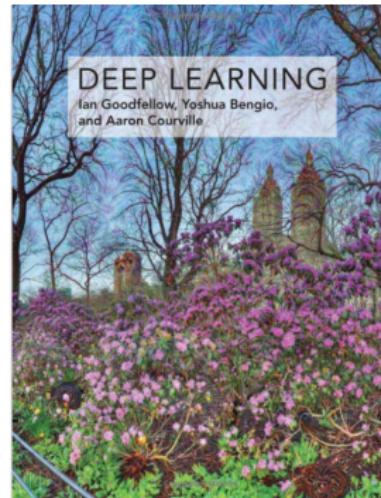
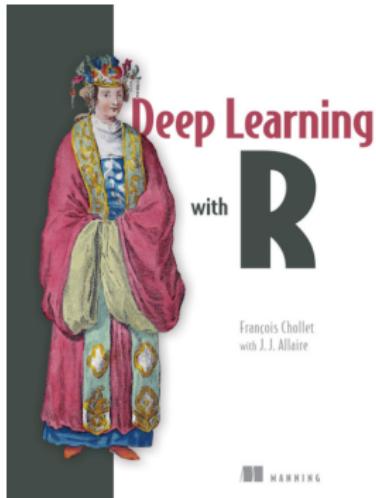
- Deep learning has made great progress and will likely increase in importance in various fields in the coming years.
- Learning representations from data that puts an emphasis on learning successive layers of increasingly meaningful representations with different types of model architectures (FFNN, CNN, RNN).
- The fundamental trick in deep learning is to use this score as a feedback signal to adjust the value of the weights a little, in a direction that will lower the loss score for the current batch of examples with SGD.
- Is not very difficult to implement nice and useful deep learning applications with the current libraries like TensorFlow and Keras.

# Deep Learning perspectives

*"Following a scientific revolution, progress generally follows a sigmoid curve: it starts with a period of fast progress, which gradually stabilizes as researchers hit hard limitations, and then further improvements become incremental. Deep learning in 2017 seems to be in the first half of that sigmoid, with much more progress to come in the next few years."*

FRANÇOIS CHOLLET

## Interesting readings and references



## Interesting readings and references

-  Andrew Ng coursera.  
<https://wwwdeeplearningai/>
-  Deep Learning - The Straight Dope.  
<http://gluon.mxnet.io/>
-  R Interface to TensorFlow.  
<https://tensorflow.rstudio.com/>
-  Keras Examples.  
[https://tensorflow.rstu](https://tensorflow.rstudio.com/keras/articles/examples/)dio.com/keras/articles/examples/

Any sufficiently advanced  
technology is  
indistinguishable from magic.

Arthur C. Clarke

e-mail: rborras@clinic.cat